

MDPs: MARKOV DECISION PROCESSES

JULIA ACADEMY: POMDPs.JL

DECISION MAKING UNDER UNCERTAINTY

WHAT IS AN MDP?

Definition: MDP. A *Markov decision process* (MDP) is a *problem formulation* that defines how an agent takes sequential *actions* from *states* in its environment, guided by *rewards*—using uncertainty in how it *transitions* from state to state.

- Formally, an MDP is defined by the following:


Table: MDP Problem Formulation: $\langle \mathcal{S}, \mathcal{A}, T, R, \gamma \rangle$

Variable	Description	POMDPs Interface
\mathcal{S}	State space	POMDPs.states
\mathcal{A}	Action space	POMDPs.actions
$T(s' \mid s, a)$	Transition function	POMDPs.transition
$R(s, a)$	Reward function	POMDPs.reward
$\gamma \in [0, 1]$	Discount factor	POMDPs.discount

Remember, an MDP is a *problem formulation* and *not an algorithm*.
An MDP formulation enables the use of solution methods, i.e. algorithms.

MDP EXAMPLE: GRID WORLD

In the **Grid World** problem, an *agent* moves around a grid attempting to collect as much reward (**green cells**) as possible, avoiding negative rewards (**red cells**).



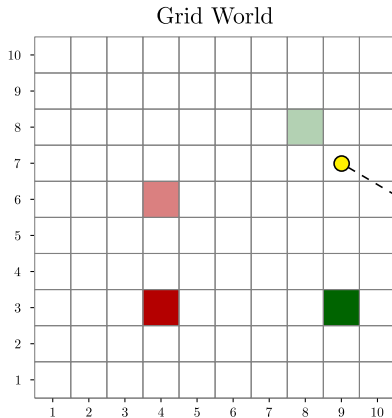
MDP: STATE SPACE

Definition: State space \mathcal{S} .

A set of all possible *states* an agent can be in (discrete or continuous).

Grid World example:

All possible (x, y)
cells in a 10×10 grid
(i.e., 100 discrete states)



State
 (x, y) of $(9, 7)$


MDP: ACTION SPACE

Definition: Action space \mathcal{A} .

A set of all possible *actions* an agent can take (discrete or continuous).

Grid World example:

The four (discrete)
cardinal directions:
[up, down, left, right]



MDP: TRANSITION FUNCTION

Definition: Transition function¹ $T(s' | s, a)$.

Defines how the agent *transitions* from the current state s to the next state s' when taking action a .
Returns a *probability distribution* over all possible next states s' given (s, a) .

Grid World example:

Stochastic transitions (incorporates randomness/uncertainty).

Action a = up from state s .

70% chance of transitioning correctly.

30% chance ($10\% \times 3$) of transitioning incorrectly.²

	0.7	
0.1	s \uparrow a	0.1
	0.1	

¹Sometimes called the *transition model*.

²i.e., a different action is taken.

MDP: REWARD FUNCTION


Definition: Reward function¹ $R(s, a)$.

A defines the *reward* an agent receives when taking action a from state s .

Grid World example:

Two cells contain **positive rewards**
and two cells contain **negative rewards**,
all others are zero.

Grid World Rewards




¹Sometimes called the *reward model*.


MDP: DISCOUNT FACTOR

Definition: Discount factor $\gamma \in [0, 1]$.


The *discount factor* controls how myopic (short-sighted) the agent is in its decision making (e.g., when $\gamma = 0$, the agent only cares about immediate rewards (myopic) and as $\gamma \rightarrow 1$, the agent takes in potential future information in its decision making process).




(a) Short-sighted
(no reward spread)



(b) Some future
reward¹ is spread



(c) Future reward
is nicely spread



(d) Dominated by
the future reward

¹The sum of the *discounted future rewards* is called the *utility* $U(s)$ or the *value* $V(s)$ of a state.

QuickPOMDPs: GRID WORLD

```
using POMDPs, POMDPModelTools, QuickPOMDPs

struct State; x::Int; y::Int end # State definition
@enum Action UP DOWN LEFT RIGHT # Action definition

s = [[State(x,y) for x=1:10, y=1:10]..., State(-1,-1)] # State-space
A = [UP, DOWN, LEFT, RIGHT] # Action-space

const MOVEMENTS = Dict{UP=>State(0,1), DOWN=>State(0,-1), LEFT=>State(-1,0), RIGHT=>State(1,0)}
Base.:+(s1::State, s2::State) = State(s1.x + s2.x, s1.y + s2.y) # Helper for applying actions


function T(s, a) # Transition function
    R(s) != 0 && return Deterministic(State(-1,-1))
    Ns = length(s)
    next_states = Vector{State}{}(undef, Ns + 1)
    probabilities = zeros{Ns + 1}
    for (i, a') in enumerate(s)
        prob = (a' == a) ? 0.7 : (1 - 0.7) / (Ns - 1)
        destination = s + MOVEMENTS[a']
        next_states[i+1] = destination
        if 1 ≤ destination.x ≤ 10 && 1 ≤ destination.y ≤ 10
            probabilities[i+1] += prob
        end
    end
    (next_states[1], probabilities[1]) = (s, 1 - sum(probabilities))
    return SparseCat(next_states, probabilities)
end

function R(s, a=missing) # Reward function
    if s == State(4,3)
        return -10
    elseif s == State(4,6)
        return -5
    elseif s == State(9,3)
        return 10
    elseif s == State(8,8)
        return 3
    end
    return 0
end

abstract type GridWorld <: MDP{State, Action} end

mdp = QuickMDP(GridWorld,
    states = s,
    actions = A,
    transition = T,
    reward = R,
    discount = 0.95,
    isterminal = s->s==State(-1,-1));
```

- This code^a defines the entire *Grid World* problem using QuickPOMDPs.jl
 - Just a sneak-peek: we'll walk through this in detail in the Pluto notebooks



^aYes, this is self-contained—copy and paste it into a notebook or REPL!

MDP SOLVERS

A number of ways to solve MDPs are implemented in the following packages.

Table: MDP Solution Methods

Package	Online/Offline	State Spaces	Actions Spaces
DiscreteValueIteration.jl	Offline	Discrete	Discrete
LocalApproximationValueIteration.jl	Offline	Continuous	Discrete
GlobalApproximationValueIteration.jl	Offline	Continuous	Discrete
MCTS.jl*	Online	Continuous	Continuous

* Monte Carlo Tree Search.

When defining your problem, the *type* of state and action space is very important!

REINFORCEMENT LEARNING SOLVERS

Certain problems are better suited in the *reinforcement learning* (RL) domain. Several RL solvers that adhere to the `POMDPs.jl` interface are implemented in the following packages.

Table: Reinforcement Learning Solution Methods

Package	State Spaces	Actions Spaces	Algorithms Implemented
TabularTDLearning.jl	Discrete	Discrete	Q-learning, SARSA, SARSA- λ
DeepQLearning.jl	Continuous	Discrete	DQN, Double DQN, Dueling DQN, Recurrent Q-learning
Crux.jl	Continuous	Continuous	DQN, REINFORCE, PPO, A2C, DDPG, TD3, SAC, Behavior Cloning, GAIL, AdvIL, AdRIL, SQIL, ASAF

When defining your problem, the *type* of state, action, and observation space is very important!